

Problem statement:-

To predict the best model for the given Rainfall dataset based on accuracy.

Data Collection

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

In [3]:

```
df=pd.read_csv(r"C:\Users\sowmika\Downloads\rainfall.csv")
df
```

Out[3]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	O
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	38
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	19
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	18
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	22
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	26
...
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	11
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	14
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	7
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	16
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	16

4116 rows × 19 columns



Data Cleaning and Preprocessing

In [4]:

```
df.head()
```

Out[4]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	388.5	5
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	197.2	3
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	181.2	2
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	222.2	3
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	260.7	



In [5]:

```
df.tail()
```

Out[5]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	117.4	
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	145.9	
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	72.8	
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	169.3	
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	165.4	



In [7]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4116 entries, 0 to 4115
Data columns (total 19 columns):
#   Column          Non-Null Count  Dtype
---  -
0   SUBDIVISION      4116 non-null   object
1   YEAR             4116 non-null   int64
2   JAN              4112 non-null   float64
3   FEB              4113 non-null   float64
4   MAR              4110 non-null   float64
5   APR              4112 non-null   float64
6   MAY              4113 non-null   float64
7   JUN              4111 non-null   float64
8   JUL              4109 non-null   float64
9   AUG              4112 non-null   float64
10  SEP              4110 non-null   float64
11  OCT              4109 non-null   float64
12  NOV              4105 non-null   float64
13  DEC              4106 non-null   float64
14  ANNUAL           4090 non-null   float64
15  Jan-Feb          4110 non-null   float64
16  Mar-May          4107 non-null   float64
17  Jun-Sep          4106 non-null   float64
18  Oct-Dec          4103 non-null   float64
dtypes: float64(17), int64(1), object(1)
memory usage: 611.1+ KB
```

In [8]:

```
df.shape
```

Out[8]:

```
(4116, 19)
```

In [9]:

```
df.isnull().any()
```

Out[9]:

```
SUBDIVISION    False
YEAR           False
JAN            True
FEB            True
MAR            True
APR            True
MAY            True
JUN            True
JUL            True
AUG            True
SEP            True
OCT            True
NOV            True
DEC            True
ANNUAL         True
Jan-Feb        True
Mar-May        True
Jun-Sep        True
Oct-Dec        True
dtype: bool
```

In [10]:

```
df.fillna(method="ffill",inplace=True)
```

In [11]:

```
df.isnull().sum()
```

Out[11]:

```
SUBDIVISION    0
YEAR           0
JAN            0
FEB            0
MAR            0
APR            0
MAY            0
JUN            0
JUL            0
AUG            0
SEP            0
OCT            0
NOV            0
DEC            0
ANNUAL         0
Jan-Feb        0
Mar-May        0
Jun-Sep        0
Oct-Dec        0
dtype: int64
```

In [12]:

```
df['YEAR'].value_counts()
```


Out[12]:

```
YEAR
1963    36
2002    36
1976    36
1975    36
1974    36
..
1915    35
1918    35
1954    35
1955    35
1909    34
Name: count, Length: 115, dtype: int64
```

Exploratory data analysis


In [13]:

```
sns.lmplot(x='NOV',y='DEC',order=2,data=df,ci=None)
plt.show()
```



In [14]:

```
df=df[['JAN','FEB','MAR','APR','DEC']]  
sns.heatmap(df.corr(),annot=True)  
plt.show()
```




In [15]:

```
sns.pairplot(df)
```

Out[15]:

<seaborn.axisgrid.PairGrid at 0x160dad30d0>



Splitting the dataset into training data and test data

In [16]:

```
x=np.array(df['FEB']).reshape(-1,1)  
y=x=np.array(df['JAN']).reshape(-1,1)
```

In [17]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.30)
```


In [18]:

```
lrg=LinearRegression()  
lrg.fit(x_train,y_train)  
print(lrg.score(x_train,y_train))
```

1.0


In [19]:

```
y_pred=lrg.predict(x_test)  
plt.scatter(x_test,y_test,color='cyan')  
plt.plot(x_test,y_pred,color='black')  
plt.show()
```



In [20]:

```
df700=df[:][:700]  
sns.lmplot(x='FEB',y='JAN',order=2,ci=None,data=df700)  
plt.show()
```



In [21]:

```
df700.fillna(method='ffill',inplace=True)
```

In [22]:

```
x=np.array(df700['FEB']).reshape(-1,1)  
y=x*np.array(df700['JAN']).reshape(-1,1)
```

In [23]:

```
df700.dropna(inplace=True)
```


In [24]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.03)
lrg=LinearRegression()
lrg.fit(x_train,y_train)
print(lrg.score(x_test,y_test))
```

1.0

In [25]:

```
y_pred=lrg.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```



In [26]:

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
```

In [27]:

```
lrg=LinearRegression()
lrg.fit(x_train,y_train)
y_pred=lrg.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2 score:",r2)
```

R2 score: 1.0


Ridge Regression

In [28]:

```
from sklearn.linear_model import Lasso,Ridge
from sklearn.preprocessing import StandardScaler
```

In [29]:

```
plt.figure(figsize=(10,10))
sns.heatmap(df.corr(),annot=True)
plt.show()
```



In [30]:

```
features=df.columns[0:5]
target=df.columns[-5]
```

In [31]:

```
x=df[features].values
y=df[target].values
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.30,random_state=1)
print("The dimension of X_train is {}".format(x_train.shape))
print("The dimension of X_test is {}".format(x_test.shape))
```

The dimension of X_train is (2881, 5)
The dimension of X_test is (1235, 5)

In [32]:

```
lrg= LinearRegression()
#Fit model
lrg.fit(x_train, y_train)
actual = y_test
train_score_lrg = lrg.score(x_train, y_train)
test_score_lrg = lrg.score(x_test, y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lrg))
print("The test score for lr model is {}".format(test_score_lrg))
```

Linear Regression Model:

The train score for lr model is 1.0
The test score for lr model is 1.0

In [33]:

```
ridgeReg = Ridge(alpha=10)
ridgeReg.fit(x_train,y_train)
train_score_ridge = ridgeReg.score(x_train, y_train)
test_score_ridge = ridgeReg.score(x_test, y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.9999999999856335
The test score for ridge model is 0.9999999999840021

Lasso Regression

In [35]:

```
lasso= Lasso(alpha=10)
lasso.fit(x_train,y_train)
train_score_ls = lasso.score(x_train, y_train)
test_score_ls= lasso.score(x_test, y_test)
print("\nLasso Model:\n")
print("The train score for lasso model is {}".format(train_score_ls))
print("The test score for lasso model is {}".format(test_score_ls))
```

Lasso Model:

The train score for lasso model is 0.9999147271297208
The test score for lasso model is 0.9999147248375002

In [36]:

```
plt.figure(figsize=(10,10))
```

Out[36]:

<Figure size 1000x1000 with 0 Axes>

<Figure size 1000x1000 with 0 Axes>

In [37]:

```
from sklearn.linear_model import LassoCV
```

In [38]:

```
from sklearn.linear_model import RidgeCV
ridge_cv=RidgeCV(alphas =[0.0001,0.001,0.01,0.1,1,10]).fit(x_train,y_train)
print(ridge_cv.score(x_train,y_train))
print(ridge_cv.score(x_test,y_test))
```

0.9999999982836236

0.9999999986591067

Elastic Net

In [39]:

```
from sklearn.linear_model import ElasticNet
```

In [40]:

```
e=ElasticNet()  
e.fit(x_train,y_train)  
print(e.coef_)  
print(e.intercept_)  
print(e.score(x,y))
```

```
[9.99044548e-01 1.38835344e-05 4.58897515e-05 0.00000000e+00  
 0.00000000e+00]  
0.016565679683701262  
0.9999991435191248
```

In [41]:

```
y_pred_elastic=e.predict(x_train)
```

In [51]:

```
mean_sqaured_error=np.mean((y_pred_elastic-y_train)**2)  
print(mean_sqaured_error)
```

```
0.0009226812593710402
```

Conclusion:

we conclude that "Ridge model" is the best model for Rainfall Prediction dataset,because it got highest accuracy compared to other models